

Admin: Basic selection rule AI integration

Overview

RPI can use Artificial Intelligence (AI) to assist in the building of criteria within basic selection rules.

Requirements

i Azure OpenAI is a proprietary service of Microsoft Corporation and requires a separate license subscription directly from Microsoft Corporation to access.

- [Azure OpenAI](#) service
 - OpenAI chat models supported:
 - GPT-5
 - GPT-5-mini
 - GPT-5.1
 - GPT-5.2
 - text-embedding-ada-002
- [Azure Search](#) service
- Azure Blob Storage account

i For guidance on which model will best suit your needs, refer to:

- [Choose the right AI model for your workload](#)
- [Choosing the right model for your use case](#)

Vector algorithm

The vector algorithm is generated by RPI automatically along with the search index, vector search, etc. The values shown in the screenshot below are what RPI applies. There is no client-facing config for the vector algorithm; RPI sets it when it generates it via “Update AI model”, which recreates the index each run (any manual edits to the index/algorithm will be overwritten). The customer cannot configure/tune the values.

Refer to the following for more information:

- [Azure AI Search](#)
- [Foundry models](#)
- [Create a vector index](#)

The following is a sample of the vector algorithm configuration:

Vector algorithm

Algorithm name * ⓘ
vector-config-1702034669774

Kind * ⓘ
hnsw

Kind parameters

Bi-directional link count (m) ⓘ
4

efConstruction ⓘ
400

efSearch ⓘ
500

Similarity metric ⓘ
cosine

Use of Azure services in this context is subject to cost, and is determined by the selected model type as well as the complexity of a given prompt and the resulting returned attributes. Please refer to the [Azure Pricing page](#) for more information.

Configuring Redpoint AI

To enable and configure Redpoint AI, use the [RPI Helm Assistant](#).

1. Set `redpointAI.enabled` to `true`.
2. Set the following:

Setting	Description
<code>naturalLanguage.ApiBase</code>	The Azure OpenAI Service name, in this format: <code>https://[Azure OpenAI Service Name].openai.azure.com/</code>

<code>naturalLanguage.ApiVersion</code>	Azure OpenAI API version.
<code>naturalLanguage.ChatGptEngine</code>	Azure OpenAI API chat model deployment. Supported models are listed in the Overview section above. This is set to the most recently tested version by default.
<code>naturalLanguage.ChatGptTemp</code>	Set to <code>0.5</code> by default, the Azure OpenAI chat model temperature controls the randomness and creativity of the model's responses, ranging from 0.0 to 2.0: • Lower values (0.0 to 0.3): Highly focused, deterministic, and factual. • Moderate values (0.7 to 1.0): A balance of coherence and creativity, resulting in natural-sounding, diverse responses. • Higher values (1.2 to 2.0): Highly creative and unpredictable.
<code>cognitiveSearch.SearchEndpoint</code>	The endpoint for the AI Search service in this format: <code>https://[AI Search Service Name].search.windows.net</code>
<code>cognitiveSearch.VectorSearchProfile</code>	Azure AI search vector search profile.
<code>cognitiveSearch.VectorSearchConfig</code>	Azure AI search vector search config.
<code>modelStorage.EmbeddingsModel</code>	Azure embeddings model deployment.
<code>modelStorage.ModelDimensions</code>	An integer setting that specifies the embedding vector size (number of dimensions) used for model embeddings in Redpoint AI model storage.

	It should match the dimensionality produced by the embedding model you're using (e.g., 768, 1024, 1536, etc.). The default setting is <code>1536</code>.
<code>modelStorage.ContainerName</code>	Azure Storage container name.
<code>modelStorage.BlobFolder</code>	Azure Storage Blob.

⚠ Sensitive values (`ApiKey` , `SearchKey` , `ConnectionString`) are stored in the K8s Secret, not here. See Required Secret Keys in [secrets-management.md](#):
`RPI_NLP_API_KEY` , `RPI_NLP_SEARCH_KEY` ,
`RPI_NLP_MODEL_CONNECTION_STRING` .

Optionally, the following setting can be used to enable trace logging:

- `RPI__NLP__EnableTrace` [Boolean]

When set to `true` , at invocation of **SQL Database Definition > Update AI Model**, each time a call is made to the OpenAI API, the following are logged to the RPI Server Log as errors:

- JSON representation of all `RPI_NLP_` settings, with keys partially obfuscated
- Endpoint called
- Payload sent
- Response received